



LITICS

AI-POWERED LEGAL AUTOMATION

NETWORK PILOT

Route every case to the firm **most likely to win it.**

An AI agent that learns your network's referral patterns and matches incoming cases to the right firm — with a visible rationale, decline-reason capture, and outcome tracking. Built for tenant law, personal injury, and the rest of the referral economy.

— THE PROBLEM

Every firm in your network has a wheelhouse — case types, jurisdictions, signal patterns they want and don't want. Today, that knowledge lives in spreadsheets, intake notes, and someone's memory.

The cost: cases declined that should have been accepted, cases accepted that go nowhere, and referral fees walking out the door on every mismatch.

— THE SOLUTION

A platform that watches every routing decision your network makes — accepts, declines, overrides, outcomes — and gets sharper at recommending the right firm for the next case.

No black box. Every suggestion comes with a rationale and a confidence level. The user always decides; the system always learns.

— HOW IT WORKS

01 · INTAKE

Cases captured with structured signals

Case type drives the right intake questions — habitability cases ask about code violations and infestations; PI asks about hospitalization and corporate defendants. Universal fields cover documentation, timing, and vulnerability.

02 · MATCH

AI recommends a firm with rationale

Reads the case, every firm's profile, and the network's full referral history. Returns a single recommendation, the reasons behind it, and a confidence rating — never just a score.

03 · LEARN

Every decision becomes training data

Accepts, declines (with reasons), overrides, and final outcomes feed back. The system gets sharper on your network's actual patterns — not generic priors.

— THE MATH

Industry estimates put case intake at 90–120 minutes per matter, with another 20–40 minutes spent on routing — deciding which firm gets it, calling around, and re-routing after declines. For a network running 50 cases per month, that is **50–100 hours of pure routing labor** before a single case is even worked.

~30 min

AVERAGE TIME STAFF SPEND
ROUTING EACH CASE TODAY

25+ hrs

STAFF HOURS PER MONTH
AT 50 CASES — ~\$1,800 IN
LOADED LABOR

3 min

ROUTING TIME WITH THE AI
AGENT — REVIEW, ACCEPT
OR OVERRIDE

Labor savings are the floor, not the ceiling. The real money is in the cases you'd have lost — the high-tier matter that goes to the wrong firm, gets declined two weeks later, and either dies in the pipeline or settles low somewhere else. **One additional captured tenant-law matter at $\$185K \times 33\% = \$61K$ in fees** — 2.5× the platform's annual cost from a single case.

— OUTCOMES FROM THE LIVE DEMO DATASET

\$248K

LEAD POISONING · TOXIC TORT

\$620K settlement — child blood-lead elevation, peeling lead paint, prior landlord notice. Routed to specialist with documented health-damage track record.

\$145K

HABITABILITY MOLD · TRIAL WIN

\$440K jury verdict — single mom + 2 kids, refused remediation, ILED upheld. AI matched to firm with similar prior win on file.

\$61K

WRONGFUL EVICTION · OVERRIDE

\$185K settlement — elderly tenant locked out during medical absence. Operator **overrode** AI's first pick; outcome data trained the model.

— FOR AGENCY NETWORKS · PILOT OFFERING

- **Multi-tenant dashboard** with per-firm reporting and network-level analytics
- **White-label** — your brand, your domain, your support email (already supports your firm logo on the dashboard)
- **Case-type-aware intake** — tenant law, PI, family, employment, and more share one engine; signals adapt per case
- **Pilot pricing reflects co-build phase** — let's scope to your network size on the call

— NEXT STEP

A 30-day pilot. We co-design the agency dashboard against your real intake flow, run live cases through the matcher, and share the uplift math at day 30 — routing time saved, decline rate before vs. after, fees captured. If it works, we scope a network-wide rollout. If it doesn't, you've lost a month, not a year.